

BeTree: Project Report

Measure where the city is dying, pay people to fix it, prove it worked.

Team **be tree** · HackXplore 2026 · Karlsruhe Citizen app `betree.me` · API `api.betree.me`

1. Executive summary

German cities are now **legally required** (Bundes-Klimaanpassungsgesetz) to measure and act on heat and drought, but they have **zero per-tree data**. Karlsruhe alone has tens of thousands of street trees, watered on fixed truck routes. That watering is slow, labour-intensive and expensive, and without per-tree data the routes inevitably miss some of the trees that are actually drying out, so those trees die. A young tree that dies costs **€5,000–15,000** to replace.

BeTree closes two gaps at once:

1. **Real per-tree data** from a cheap solar ESP32 mesh sensor: soil moisture, temperature, an activity index and noise from one €24 build.
2. **A verified citizen action loop**: people find a thirsty tree in the app, water it, and the *sensor confirms a real moisture rise* (rain cross-checked). No honor system. Verified care earns credits redeemable for real city perks.

The economics are decisive: **one saved tree pays for a full season of citizen rewards roughly 10–30× over**, and the same sensor answers questions for **five city departments** from one install. The city buys the climate-compliance data it's legally mandated to produce anyway, and citizen labour does the saves where it pays off most: the young trees that die fastest and that volunteers can actually help.

This is a working product, not a concept

Proof	What exists today
Hardware mesh	ESP32 dual-node ESP-NOW mesh firmware (PlatformIO), low-power sensor node + WiFi gateway
Self-built prototype	We 3D-designed & printed the " Sprig " enclosure and assembled it with off-the-shelf sensors, some bought Saturday morning in a Karlsruhe shop . It's live in a tree
Live data	60,000+ real readings from two live nodes : one mesh gateway, one low-power sensor node (KA-00001 / KA-00002)
Real city	126,434 actual Karlsruhe trees (open cadastre) on the live map
Real model	Hybrid FAO-56 water balance fusing Open-Meteo/DWD weather with our live sensor data , reading root-zone depth
Shipped	Deployed PWA + FastAPI, green CI (ruff · mypy · pytest, Python 3.11–3.13)

2. The problem

- **The mandate exists, the visibility doesn't.** Cities must report and act on climate risk, but manage trees with clipboards and weather-model guesses. Nobody knows which individual tree is dying.
- **Watering is laborious, expensive, and blind.** Fixed truck routes cost time and money, and with no per-tree data they miss some of the trees that are actually drying out. Establishment-phase trees (age ≤ 5) die fastest and cost the most to replace.
- **Existing citizen apps run on trust.** "I watered this tree": self-reported, unverifiable, gameable. An expensive, city-scale problem running on a paper receipt.

3. The solution

A three-sided platform on one shared data layer:

- **Citizens** get a mobile PWA: a live map of every tree's thirst, turn-by-turn to the nearest rescue, QR/NFC tap-to-water, a live moisture chart, confetti on verified success, and a reward shop.
- **The city** gets a dashboard: health overview, activity-index and heat heatmaps, watering-route hints, and an exportable audit trail of every verified watering.
- **Five departments** get continuous street-level data (water, heat, activity index, noise, storm tilt) as a *side effect* of the tree network.

4. What we built this weekend (proof of realness)

We deliberately built **down the whole stack**, not just a pretty front end:

- **Hardware.** Designed the **Sprig** sensor housing in CAD (`3d/SprigV2.step`), 3D-printed it, and assembled it with a capacitive soil-moisture probe, DHT11, vibration sensor and microphone. Some of the parts we bought **Saturday morning at a shop in Karlsruhe**. The sensors are standard, reusable, and quick to mount on a tree. It's staked in a live tree, powered, and reporting.
- **Firmware.** A real **dual-node ESP-NOW mesh**: low-power sensor nodes deep-sleep and relay JSON over ESP-NOW to a WiFi gateway that POSTs to the backend. Each node declares which sensors it has; the backend accepts partial payloads.
- **Backend.** FastAPI + async SQLAlchemy, sensor ingest, tree/forecast/rescue endpoints, gamification, rewards, city dashboard, with **60k+ real readings** in the DB and the **126k-tree** Karlsruhe cadastre on the map.
- **Model.** A **hybrid** model: a FAO-56 soil-water balance (a standard agronomic water-accounting method) that fuses live Open-Meteo / DWD weather with our own sensor readings.
- **App.** A deployed, installable PWA with the full citizen + city experience.



The Sprig prototype: 3D-printed housing, off-the-shelf sensors, an ESP32, running our mesh firmware.

5. How the technology works

5.1 One cheap node, five data streams

Sensor	Measures	City value
Capacitive soil moisture	Root-zone hydration	Which trees need water now?
DHT11 temp + humidity	Per-tree microclimate	Urban heat-island mapping
Vibration / shock	Activity index + storm tilt	Pedestrian activity • storm-damage alerts
Microphone	Ambient sound level	Real-time noise map
Solar + LiPo	Self-powered	Deploy-and-forget

~€24 bill of materials at volume. Nodes mesh over **ESP-NOW** at zero per-node connectivity cost; a single gateway bridges a cluster to WiFi. LoRaWAN (5 km range, existing TTN infrastructure) is a radio swap on the roadmap.

5.2 Seeing into the depths, and covering the city with the right hardware

A capacitive probe only reads the **top few centimetres** of soil. But a tree drinks from the **root zone, 20–30 cm down**, which can be bone-dry while the surface looks fine. So we don't trust the surface reading alone. We run a **FAO-56 soil-water balance** (the UN Food and Agriculture Organization's standard agronomic method for tracking, day by day, how much water a plant's root zone gains and loses) driven by free weather from **Open-Meteo** (an open weather API) and **DWD ICON-D2** (the German national weather service's 2 km forecast grid over Germany):

- **Loss** = reference evapotranspiration (ET_0) × crop coefficient (per species) × a water-stress factor
- **Gain** = effective rainfall (canopy throughfall) + known citizen waterings
- **Depth** = the model explicitly pulls Open-Meteo's modelled **deep-soil moisture (3–9 cm, 9–27 cm)** as a physical prior, so it reasons about the root zone, not just the skin of the soil.

The model is **hybrid: it fuses our live sensor data with weather physics**, anchored to a real sensor where one exists and inferred from physics everywhere else. That is the key to scale: **a few hundred calibration sensors ground-truth a physical model that then predicts water stress for all 126,000 trees**, sensed and unsensed alike. It is also how we decide where the hardware actually goes:

- **Old, established trees: sparse sensors, modelled.** A mature tree has deep roots and rarely needs rescuing, so a handful of calibration sensors plus the weather model cover whole districts of them.
- **Young, establishment-phase trees: dense sensors.** These are the trees that die without help and cost the most to replace, and they are exactly where citizen watering pays off. So we instrument them densely and point people at them.

Sparse hardware for citywide coverage, dense hardware where lives are actually on the line.

We present the model as "**physics-based, calibrates as sensors deploy**": the constants are honest physical estimates, not a trained accuracy number we can't yet back.

5.3 Proof of Care: sensor-verified, tamper-resistant

Proof of Care matters most for the young trees citizens can actually help, which is exactly where the dense sensors sit. A watering only earns credit if the **soil agrees**:

- The sensor must show a **sustained moisture rise** during the session; a single spike doesn't count.
- The rise is **cross-checked against live rainfall**: credit is suppressed if it's raining.
- The session has a minimum window, so there's no fake tap-and-run.

The city gets **ground-truth**, not citizen diaries, which is exactly what the Klimaanpassungsgesetz needs. (Verification runs in the app today against the live sensor; hardening it server-side is the next step.)

6. One sensor, five departments

The network isn't a tree sensor; it's the **first street-level data grid** for the city. One install, one API, five paying stakeholders:

Department	Question the same sensor answers
Grünflächenamt	Which of our trees need water right now?
Stadtplanung	Did pedestrianising this street change foot traffic?
Umweltamt	Did Tempo-30 reduce noise on residential streets?
Gesundheitsamt	Where are heat islands exceeding 38 °C in a heat wave?
Tiefbauamt	Which trees are leaning after last night's storm?

Each department currently pays for one-off studies and manual surveys to answer these. BeTree gives them **continuous, per-street answers**, and each co-funds a fraction, so the sensor pays for itself before a single tree is saved.

7. Costs & economics

7.1 Per-node cost

Item	Approx. cost (at volume)
ESP32 + capacitive moisture + DHT11 + vibration + mic	~€24 BOM
Connectivity (ESP-NOW mesh)	€0 per node (shared gateway)
Power (solar + LiPo)	self-powered, no grid, ~zero maintenance

7.2 The order-of-magnitude argument

	Cost
Replacing one dead young tree (removal + replant + 3–5 yr establishment watering)	€5,000–15,000
A full summer of citizen watering rewards to keep a tree alive (redeemable perks)	up to ~€500
Return on a single saved tree	~10–30×

Reward catalogue (low marginal cost to the city): seed packets · museum entry · transit day-pass · priority Bürgeramt appointment. The reward is mostly **existing city capacity**, not cash out the door.

7.3 Pilot budget (illustrative)

A **50-node, one-district, one-summer** pilot: ~€1,200 in hardware + a gateway, the app and API already built, weather data free. If the network keeps **even one or two trees alive that would otherwise die**, it has already **paid for itself**, before counting the cross-department data value.

8. Scalability & realisability

Scales like software, sells like infrastructure.

- **Coverage scales with the model, not the hardware.** Because the physical model generalises from a few hundred calibration nodes to the whole cadastre, you don't need a sensor per tree. Coverage of 126k trees is a software property; accuracy tightens as you add sensors.
- **Connectivity scales cheaply.** ESP-NOW mesh is free per node; LoRaWAN (roadmap) reuses existing TTN gateways for km-scale range with no per-tree WiFi.
- **Data scales the customer base.** One install serves five departments, so the addressable value grows without more hardware.
- **The software is production-shaped already.** Async FastAPI, typed schemas, CI on three Python versions, a clean migration path SQLite → Postgres/PostGIS for spatial queries.

Realisability check: what's real vs. roadmap today:

Real now	Roadmap
ESP-NOW mesh firmware, 3D-printed node, live sensor	LoRaWAN radio swap
60k+ readings, 126k-tree map, deployed app + API	Server-side verification hardening (anti-gaming)
FAO-56 + Open-Meteo model, forecast endpoint	Calibrate constants on real sensor history; ML residual layer
CI, tests, typed code	Postgres + PostGIS, push notifications

9. Technical feasibility

- **Frontend:** Next.js 16, React 19, TypeScript, Tailwind 4, MapLibre GL vector tiles, Zustand, PWA.
- **Backend:** FastAPI, async SQLAlchemy 2, SQLite (→ Postgres/PostGIS), Pydantic 2, Open-Meteo client.
- **Firmware:** ESP32 (Arduino/PlatformIO), ESP-NOW mesh, modular per-node sensor config.
- **Quality:** ruff + mypy + pytest in GitHub Actions across Python 3.11/3.12/3.13.
- **Model:** pure, unit-tested water-balance functions; free weather data, cached per neighbourhood.

Every claim above is backed by code in this repository; see the [root README](#).

10. Risks & honest limitations

- **Model not yet calibrated.** Constants are physics-shaped estimates; we deliberately quote no accuracy figure until we fit on real dry-down data.
- **Verification is app-side today.** Anti-gaming logic exists in the flow but must move server-side to be tamper-proof at scale; it's the first roadmap item.
- **Two live sensors** (one mesh gateway, one low-power node). The pipeline is real end-to-end; statistical confidence across the cadastre needs more nodes.

BeTree: because the trees can't ask for help themselves, but your phone can.